

Estimating Models by The Generalized Method of Moments

Elad Guttman

Outline

1. Working with panel data

2. GMM

2.1 A Reminder

2.2 Examples

Working with panel data

dplyr basic verbs

- `mutate()` - adds new variables that are functions of existing variables
- `select()` - picks variables based on their names.
- `filter()` - picks cases based on their values.
- `summarise()` - reduces multiple values down to a single summary.
- `arrange()` - changes the ordering of the rows.

Leads and Lags

Another two useful dplyr functions for working with panel data are the lead and lag functions:

- lag - Find the "previous" values in a vector.
- lead - Find the "next" values in a vector.

```
library(tidyverse)
x = 1:10
lag(x, n = 1)
```

```
## [1] NA  1  2  3  4  5  6  7  8  9
```

```
lag(x, n = 2)
```

```
## [1] NA NA  1  2  3  4  5  6  7  8
```

```
lead(x, n = 1)
```

```
## [1]  2  3  4  5  6  7  8  9 10 NA
```

Combining all together

Let's combine it all together to calculate the daily log return ($r_t = \log(\frac{p_t}{p_{t-1}})$) on several stock prices:

```
library(tidyquant)
stocks = c("AAPL", "NFLX", "AMZN")
prices = tq_get(stocks,
                from = "2020-01-01",
                to = "2020-03-01",
                get = "stock.prices")
head(prices, 4)
```

```
## # A tibble: 4 x 8
##   symbol date      open  high  low close  volume adjusted
##   <chr> <date>    <dbl> <dbl> <dbl> <dbl>    <dbl>    <dbl>
## 1 AAPL  2020-01-02  74.1  75.2  73.8  75.1  135480400  74.4
## 2 AAPL  2020-01-03  74.3  75.1  74.1  74.4  146322800  73.7
## 3 AAPL  2020-01-06  73.4  75.0  73.2  74.9  118387200  74.3
## 4 AAPL  2020-01-07  75.0  75.2  74.4  74.6  108872000  74.0
```

Combining all together

```
prices = prices %>%  
  arrange(symbol, date) %>%  
  group_by(symbol) %>%  
  mutate(log_return = log(open/lag(open))) %>%  
  ungroup()
```

```
head(prices %>% filter(symbol == "AAPL"), 3)
```

```
## # A tibble: 3 x 9
```

##	symbol	date	open	high	low	close	volume	adjusted	log_return
##	<chr>	<date>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
## 1	AAPL	2020-01-02	74.1	75.2	73.8	75.1	135480400	74.4	NA
## 2	AAPL	2020-01-03	74.3	75.1	74.1	74.4	146322800	73.7	0.00307
## 3	AAPL	2020-01-06	73.4	75.0	73.2	74.9	118387200	74.3	-0.0114

```
head(prices %>% filter(symbol == "NFLX"), 3)
```

```
## # A tibble: 3 x 9
```

##	symbol	date	open	high	low	close	volume	adjusted	log_return
##	<chr>	<date>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
## 1	NFLX	2020-01-02	326.	330.	325.	330.	4485800	330.	NA
## 2	NFLX	2020-01-03	327.	330.	326.	326.	3806900	326.	0.00208
## 3	NFLX	2020-01-06	323.	336.	321.	336.	5663100	336.	-0.0113

GMM

A Reminder

The Estimator

- We defined the objective function:

$$J_N(\theta) = \left[\frac{1}{N} \sum g(w_i, \theta) \right]' W \left[\frac{1}{N} \sum g(w_i, \theta) \right]$$

- Then we defined the GMM estimator as:

$$\hat{\theta}_{GMM} = \operatorname{argmin} J_N(\theta)$$

- For every weighting matrix W we get a consistent estimator for θ , but the following is the optimal one:

$$\hat{W}_{opt} = \hat{\Lambda}^{-1} \text{ where } \hat{\Lambda} = \frac{1}{N} \sum [g(w_i, \hat{\theta}) g(w_i, \hat{\theta})']$$

The Estimator

- Initially, we can't calculate \hat{W}_{opt} since we don't know $\hat{\theta}$
- The solution is to estimate GMM in two steps:
 1. Estimate $\hat{\theta}_{step1}$ using the identity matrix (or any other matrix) as a weighting matrix.
 2. Estimate $\hat{\theta}_{GMM}$ using $\hat{\Lambda}_{step1}^{-1}$ as a weighting matrix, to get the efficient **Two-Step GMM Estimator**
- Another option is to keep iterating until convergence is obtained (the **Iterated GMM Estimator**)

Examples

The `momentfit` package

- `momentfit` is the R package for estimating models by GMM
- The package supports 3 different ways (or classes) to represent the moment conditions:
 1. The “`linearModel`” class - for the special case of linear models
 2. The “`formulaModel`” class - for the special case of Minimum Distance Estimation
 3. The “`functionModel`” class - for the general case

- We'll use the `bwght` dataset, where we consider the following model:

$$\log(bwght) = \beta_0 + \beta_1 male + \beta_2 parity + \beta_3 \log(faminc) + \beta_4 packs + u$$

to estimate the effect of cigarette smoking on the weight of newborns, using cigarette price as an instrument

- In PS9 you saw that in fact cigarette price is not a good instrument, so we'll not try to interpret the results

The “linearModel” class

```
library(wooldridge)
library(momentfit)
library(lfe)

data("bwght")

linModel = momentModel(g = lbwght ~ male + parity + lfaminc + packs,
                      x = ~ male + parity + lfaminc + cigprice,
                      data = bwght,
                      vcov = "iid") #vcov under homoskedasticity

linModel
```

```
## Model based on moment conditions
## *****
## Moment type: linear
## Covariance matrix: iid
## Number of regressors: 5
## Number of moment conditions: 5
## Number of Endogenous Variables: 1
## Sample size: 1388
```

The “linearModel” class

```
gmm = gmmFit(linModel, type = "twostep")  
  
#for comparison:  
iv = felm(lbwght ~ male + parity + lfaminc | 0 | (packs ~ cigprice), data = bwght)
```

What would happen if we set `type = "onestep"` instead?

The Just-Identified Case

- Remember that in case when g is linear, i.e., $g(w_i, \theta) = z_i'(y_i - x_i\theta)$, we have a closed formula for $\hat{\theta}_{GMM}$:

$$\hat{\theta}_{GMM} = ((X'Z) \cdot W \cdot (Z'X))^{-1} \cdot (X'Z) \cdot W \cdot Z'y$$

- When $K = L$ (the just-identified case), $X'Z$ and W are square matrices, and therefore the GMM estimator reduces to:

$$\hat{\theta}_{GMM} = (Z'X)^{-1} Z'y$$

- So in this case W plays no role

The Just-Identified Case

```
summary(iv)$coefficients
```

##	Estimate	Std. Error	t value	Pr(> t)
## (Intercept)	4.467861478	0.25882893	17.26183221	1.286386e-60
## male	0.029820508	0.01777901	1.67728754	9.371224e-02
## parity	-0.001239075	0.02193217	-0.05649578	9.549550e-01
## lfaminc	0.063645997	0.05701281	1.11634549	2.644682e-01
## `packs(fit)`	0.797106270	1.08627520	0.73379772	4.631964e-01

```
summary(gmm)@coef
```

##	Estimate	Std. Error	t value	Pr(> t)
## (Intercept)	4.467861478	0.25845543	17.28677707	5.916934e-67
## male	0.029820508	0.01775335	1.67971137	9.301349e-02
## parity	-0.001239075	0.02190052	-0.05657742	9.548818e-01
## lfaminc	0.063645997	0.05693054	1.11795871	2.635846e-01
## packs	0.797106270	1.08470771	0.73485812	4.624259e-01

The “functionModel” class

We can estimate the same model by defining g directly:

```
g = function(theta, dat){  
  #extract the relevant variables:  
  X = dat %>%  
    select(intercept, male, parity, lfaminc, packs) %>%  
    as.matrix()  
  Z = dat %>%  
    select(intercept, male, parity, lfaminc, cigprice) %>%  
    as.matrix()  
  y = matrix(dat$lbwght, ncol = 1)  
  beta = as.matrix(theta, ncol = 1)  
  u = as.vector(y - X%*%beta)  
  return(Z*u)  
}  
  
#add an intercept  
bwght = bwght %>% mutate(intercept = 1)  
#initial values for the numerical algorithm:  
theta0 = rnorm(5)  
names(theta0) = c("intercept", "male", "parity", "lfaminc", "packs")  
funModel = momentModel(g = g,  
                        x = bwght,  
                        theta0 = theta0,  
                        vcov = "iid") #vcov under heteroskedasticity
```

The “functionModel” class

```
funModel
```

```
## Model based on moment conditions
## *****
## Moment type: function
## Covariance matrix: iid
## Number of regressors: 5
## Number of moment conditions: 5
## Number of Endogenous Variables: 0
## Sample size: 1388
```

The “functionModel” class

```
gmm2 = gmmFit(funModel, type = "twostep")
summary(iv, robust = T)$coefficients
```

##	Estimate	Robust s.e	t value	Pr(> t)
## (Intercept)	4.467861478	0.25631403	17.43120118	1.142829e-61
## male	0.029820508	0.01722088	1.73164842	8.355917e-02
## parity	-0.001239075	0.02537546	-0.04882966	9.610621e-01
## lfaminc	0.063645997	0.05707269	1.11517428	2.649695e-01
## `packs(fit)`	0.797106270	1.11322077	0.71603611	4.740899e-01

```
summary(gmm2)@coef
```

##	Estimate	Std. Error	t value	Pr(> t)
## intercept	4.467861773	0.25585173	17.46269922	2.755897e-68
## male	0.029820502	0.01718982	1.73477673	8.278036e-02
## parity	-0.001239055	0.02532974	-0.04891702	9.609854e-01
## lfaminc	0.063645932	0.05696979	1.11718742	2.639142e-01
## packs	0.797105077	1.11121370	0.71732834	4.731715e-01

The Over-Identified Case

- Now things are getting more interesting...
- We'll estimate the model using 2 different weighting matrices: the 2SLS matrix and the optimal matrix (Do we expect to get different results?)
- But first need to modify g :

```
g = function(theta, dat){  
  #extract the relevant variables:  
  X = dat %>%  
    select(intercept, male, parity, lfaminc, packs) %>%  
    as.matrix()  
  Z = dat %>%  
    select(intercept, male, parity, lfaminc, cigprice, cigprice2) %>%  
    as.matrix()  
  y = matrix(dat$lbwght, ncol = 1)  
  beta = as.matrix(theta, ncol = 1)  
  u = as.vector(y - X*beta)  
  return(Z*u)  
}
```

The Over-Identified Case

```
#generate more instrument
bwght = bwght %>% mutate(cigprice2 = cigprice^2)
model = momentModel(g = g,
                    x = bwght,
                    theta0 = theta0,
                    vcov = "iid")

#estimate with the optimal matrix
res_with_optimal_mat = gmmFit(model, type = "twostep")

#estimate with the 2sls matrix

#define (Z'Z)^-1
Z = bwght %>%
  select(intercept, male, parity, lfaminc, cigprice, cigprice2) %>%
  as.matrix()
W = solve(t(Z)%*%Z)
res_with_2sls_mat = gmmFit(model, type = "onestep", weights = W)
```

The Over-Identified Case

```
summary(res_with_optimal_mat)@coef
```

```
##           Estimate Std. Error   t value    Pr(>|t|)
## intercept  4.457063590 0.21032171 21.1916476 1.140364e-99
## male       0.029931177 0.01775034  1.6862310 9.175132e-02
## parity    -0.002423397 0.01942087 -0.1247831 9.006952e-01
## lfaminc    0.066078306 0.04652062  1.4204089 1.554887e-01
## packs      0.846892950 0.87865412  0.9638525 3.351199e-01
```

```
summary(res_with_optimal_mat)@specTest
```

```
##
## J-Test
##           Statistics  df  pvalue
## Test E(g)=0:    0.0044147  1  0.94702
```

The Over-Identified Case

```
summary(res_with_2sls_mat)@coef
```

##		Estimate	Std. Error	t value	Pr(> t)
##	intercept	4.458513135	0.20885780	21.34712300	4.146951e-101
##	male	0.029986374	0.01765379	1.69857987	8.939837e-02
##	parity	-0.001966193	0.02006417	-0.09799525	9.219361e-01
##	lfaminc	0.065698507	0.04621798	1.42149246	1.551736e-01
##	packs	0.836839472	0.87549936	0.95584248	3.391518e-01

```
summary(res_with_2sls_mat)@specTest
```

##		Statistics	df	pvalue
##	J-Test			
##	Test E(g)=0:	0.0053864	1	0.94149

Minimum Distance Estimation

- This is a special case when moments have data separately from parameters
- For example, consider the problem of estimating the parameters of the normal distribution μ and σ^2
- The moments conditions are:
 1. $x_i - \mu$
 2. $x_i^2 - \mu^2 - \sigma^2$
- Let's go back to the prices data and estimate these parameters for the distribution of log-returns (which in this specific case seems like the normal distribution)

The “formulaModel” class

```
moment_conditions = list(log_return ~ mu,  
                          log_return^2 ~ mu^2 + sigma2)  
  
theta0 = c(0.1, 0.1)  
names(theta0) = c("mu", "sigma2")  
formulaModel = momentModel(g = moment_conditions,  
                           theta0 = theta0,  
                           data = prices,  
                           vcov = "CL", #clustered SE  
                           #also need to provide the clustering variable:  
                           vcovOptions = list(cluster = ~ symbol))  
  
formulaModel
```

```
## Model based on moment conditions  
## *****  
## Moment type: formula  
## Covariance matrix: CL  
## Clustered based on: symbol  
## Number of regressors: 2  
## Number of moment conditions: 2  
## Number of Endogenous Variables: 0  
## Sample size: 117
```

The “formulaModel” class

```
res = gmmFit(formulaModel, type = "twostep")
summary(res)@coef
```

```
##           Estimate  Std. Error   t value   Pr(>|t|)
## mu      -0.0005408820 1.868354e-03 -0.2894965 7.722014e-01
## sigma2  0.0006429251 5.811574e-05 11.0628390 1.899870e-28
```

```
#convergence status for the numerical optimization algorithm:
summary(res)@convergence
```

```
## [1] 0
```